



# Linear Algebra

## on a distributed environment

---

Jing Liu

TDB & LMB, Uppsala University

Programming of Parallel Computers, Jan, 2015



# Topics in LA

- Solve linear equations
- Matrices operations
  - ✱ Matrices addition/ multiplication / transformation
  - ✱ Eigenvalue/ Eigenvector
  - ✱ Transpose, projection ...
- Vector space
- ...

Scalar  
Vector  
Matrix



# Loops

- LA operations are often basic building blocks in scientific applications
- Three basic types of loops
  - ✱ Perfectly parallel loops
  - ✱ Reduction loops
  - ✱ Recursive loops
  - ✱ Combination of different loops



# Perfectly parallel loops

- Example  $Z_m = \lambda X_m + Y_m$   
for (  $i = 0$ ;  $i < m$ ;  $i++$  ) {  
     $Z[i] = \lambda * X[i] + Y[i]$ ;  
}

X,Y,Z




- MPI Scatter and MPI Gather



# Reduction loops

- Limited parallelism
- Example: Dot production  $s = X \bullet Y^T$   
for (  $i = 0$ ;  $i < m$ ;  $i++$  ) {  
     $s += X[i] + Y[i]$ ;  
}



- MPI Reduce, MPI Allreduce



# Recursive loops

- Each iteration depends on the previous one
- Hardly parallelize, “serial” loop
- Example

```
for ( i = 1; i < m; i ++ ) {  
    X[ i ] = X[ i ] + X[ i - 1 ];  
}
```



# Nested loops

- Often the order of loops can be interchanged → for maximal parallelism, choose the perfectly-parallel loops as outmost, and parallelize over it.

- Example: Matrix-Vector multiplication

```
for ( i = 0; i < m; i++){  
    for ( j = 0; j < m ; j++){  
        Y[i] += A[i][j] * X[j];  
    }  
}
```



# Nested loops – Alt 1

- Row-wise partition



- All processors have a copy of  $X$ , one piece of  $A$  and  $Y$ .





# Nested loop – Alt. 2

- Block algorithm with 1D partition

$$\begin{bmatrix} Y_1 \\ Y_2 \\ Y_3 \\ Y_4 \end{bmatrix} = \begin{bmatrix} A_{00} & A_{01} & A_{02} & A_{03} \\ A_{10} & A_{11} & A_{12} & A_{13} \\ A_{20} & A_{21} & A_{22} & A_{23} \\ A_{30} & A_{31} & A_{32} & A_{33} \end{bmatrix} * \begin{bmatrix} X_0 \\ X_1 \\ X_2 \\ X_3 \end{bmatrix}$$

- Step 1: Compute  $Y[i] = A[i][i] * X[i]$  in process  $i$ , and then shift  $X[i]$  circular one step up.
- Step 2: Compute again, in which  $j=(i+1) \bmod p$ , shift  $X$  circular one step up.
- Repeat, in total  $(p-1)$  step



# Nested loop – Alt. 2 cont.

- Non-blocking communication to shift X, before computation. MPI\_Isend, MPI\_Irecv, MPI\_wait
- Which one is more efficient?
  - ✱ Alt. 2 is more memory efficient.
  - ✱ CPU efficient is all depends on the problem size, computer systems, implementations of MPI\_functions, etc.



# Nested loop – Alt. 3

- Block algorithm – 2D partition
- Processor block  $\sqrt{p} * \sqrt{p}$ ,
- Step 1: Divide  $A_{mn}$  to  $\sqrt{p} * \sqrt{p}$  blocks,  $X$  to  $\sqrt{p}$  parts
- Step 2: Processor  $P_{ij}$  get block  $A_{ij}$  and  $X_j$ , and hold  $Y_i^{(j)} = \mathbf{0}$
- Step 3:  $P_{ij}$  computes  $Y_i^{(j)} = A_{ij} * X_j$
- Step 4: Accumulate  $Y_i$  in each row.




# Nested loop – Alt. 3

$Y_0$	$\leftarrow \sum_j$	$Y_0^0 = A_{00} * X_0$	$Y_0^1 = A_{01} * X_1$	$Y_0^2 = A_{02} * X_2$	$Y_0^3 = A_{03} * X_3$
$Y_1$	$\leftarrow \sum_j$	$Y_1^0 = A_{01} * X_0$	...	...	...
$Y_2$	$\leftarrow \sum_j$	$Y_2^0 = A_{02} * X_0$	...	...	...
$Y_3$	$\leftarrow \sum_j$	$Y_3^0 = A_{03} * X_0$	$Y_3^1 = A_{13} * X_1$	$Y_3^2 = A_{23} * X_2$	$Y_3^3 = A_{33} * X_3$

- Efficient for large matrices.
- Scalability? 2D > 1D. For many processors, 1D partition strips become so thin and communications increases faster.



# Matrix-Matrix Multi.



$$C(i,j) = \sum_k A(i,k) B(k,j)$$



# More nested Loops

- Example : Matrix-Matrix Multiplication
- i and j are perfectly parallel loops, k is reduction loop

```
for ( _ = 0; _ < m; _++){  
    for ( _ = 0; _ < m ; _++){  
        for ( _ = 0 ; _ < m; _ ++){  
            C[i][j] += A[i][k] * B[k][j];  
        }  
    }  
}
```



# Matrix-Matrix Multi.

- 1D partitioning – choose  $j$  as the outmost loop  $\rightarrow$  partition data column wise

$C_0$	$C_1$	$C_2$	$C_3$
-------	-------	-------	-------

 $=$ 

$A_0$	$A_1$	$A_2$	$A_3$
-------	-------	-------	-------

 $*$ 

$B_{00}$	$B_{01}$	$B_{02}$	$B_{03}$
$B_{10}$	...		
$B_{20}$		...	
$B_{30}$			$B_{33}$

- $\rightarrow C_0 = A_0 * B_{00} + A_1 * B_{10} + A_2 * B_{20} + A_3 * B_{30}$

- ...



# $C = A * B$ , 1D partition

- A is needed in every processor.
- Alt. 1 : Every processor has completed A,  
→ Not scalable (memory?!)
- Alt. 2: Shift A around.
  - Similar idea to matrix-vector alt. 2.
  - For many processors, the stripes (block-columns) become thin and comm. overhead becomes large.





# **C = A\*B, 2D partition**

- Choose both i and j outmost.
- $\sqrt{p} * \sqrt{p}$  blocks, each processor gets one block of each matrix.
- In processor  $P_{ij}$ , compute  $C_{ij} = \sum_{k=0}^{\sqrt{p}-1} A_{ik} * B_{kj}$ 
  - ➔  $P_{ij}$  need all blocks  $A_{ik}$  in block row i, and  $B_{kj}$  in block column j
  - ➔ Communications needed.



# **$C = A*B$ , 2D partition, Alt. 1**

- Simple and naïve method.
- Simply distribute  $A$  in each block row, and distribute of  $B$  in each block column, using MPI\_functions
  - ➔ limited scalability due to memory.
    - ➔ Bad performance if data don't fit in cache



# $C = A * B$ , 2D partition, Alt. 2 Cannon's Algorithm (1969)

- Shift and compute.  $M * M$  mesh ( $\sqrt{p} * \sqrt{p}$  blocks processors, data).
- Phase 1: shift
  - ✱ Shift the  $i$  th block row of  $A$   $i$  steps cyclically to the left.
  - ✱ Shift the  $j$  th block column of  $B$   $j$  steps cyclically upwards

$A_{00}$	$A_{01}$	$A_{02}$	$A_{03}$
$A_{11}$	$A_{12}$	$A_{13}$	$A_{10}$
$A_{22}$	$A_{23}$	$A_{20}$	$A_{21}$
$A_{33}$	$A_{30}$	$A_{31}$	$A_{32}$

$B_{00}$	$B_{11}$	$B_{22}$	$B_{33}$
$B_{10}$	$B_{21}$	$B_{32}$	$B_{03}$
$B_{20}$	$B_{31}$	$B_{02}$	$B_{13}$
$B_{30}$	$B_{01}$	$B_{12}$	$B_{23}$



# $C = A * B$ , 2D partition, Alt. 2 Cannon's Algorithm Cont.

- Phase 2: Compute and shift
- For each iteration do:
  - ✱ Compute  $C_{ij} = A_{ik} * B_{kj}$  in each processor  $P_{ij}$ , where  $k = (i+j+1) \bmod M$ , where  $l$  is the number of iterations (start from 0).
  - ✱ Shift A one step left, B one step upwards
- In total,  $M-1$  steps. We can do shift with non-blocking communication, and compute while sending.
- Read more on-line [Cannon's algorithm](#).



# $C = A * B$ , 2D partition, Alt. 3 Fox's Algorithm

- In total  $M-1$  step.
- For each step  $k$  ( $k = 0, 1, \dots, M-1$ )
  - ✱ Broadcast block  $n$  of  $A$  within each block row  $i$  ( $n = (i+k) \bmod M$ )
  - ✱ Multiply the broadcasted block with  $B$ -block in each processor ( $C_{ij} += A_{in} * B_{nj}$ )
  - ✱ Shift blocks of  $B$ , one step upwards.



# $C = A * B$ , 2D partition


- Both Cannon's and fox's algorithm is scalable.
- Which is more efficient?
  - ✱ Depends on problem size, computer system, efficiency of MPI, etc
- Read more about efficient [on Fox and Cannon.](#)



# Advanced Topic: Recursive loop

## ■ Example:

```
for ( i=1; i<n; i++){  
    X[i] += X[i-1];  
}
```





# Assignment 1

- Dense matrix-matrix multiplication.
- Fortran/C/C++ and MPI.
- Two parameters:
  - ✱ The number of process
  - ✱ The size of matrices
- Randomly generate A and B
- Distribute data
- Implement Fox's algorithm
- Collect data and output.





# Assignment 1, cont.

- Data generation (at rank 0): `srand()`, `rand()` / `CALL RANDOM_SEED()`, `CALL RANDOM_NUMBER()`
- Data distribution: use `MPI_Type_vector`, `MPI_Cart_rank`, `MPI_Isend`, `MPI_Recv`
- Data Collection: `MPI_Probe`, `MPI_Cart_coords`, `MPI_Recv`, `MPI_wait`



# Assignment 1, cont.

- C structure / C++ class is helpful to make a nicer code.
  - ✱ Name space works for large project.
- Good [coding style](#) makes your code more understandable and maintainable.
- Write comments in your code to help yourself and others.
- Demo code at [https://github.com/JinLi971/MPI\\_DEMO](https://github.com/JinLi971/MPI_DEMO)



# More Advanced Topic: BLAS

## ■ CPUs:

- ✱ [Armadillo](#) : Matlab style, C++ coding .
- ✱ [CBLAS](#) : GNU supported.
- ✱ Support: AMD -> ACML; IBM -> ESSL;  
Apple -> Accelerate framework; HP -> MLIB;  
SUN -> Sun Performance Library,  
[Intel Math Kernel Library](#)


## ■ GPUs: NVIDIA -> [CuBLAS](#)

OPENCL : third part support.



# More Advanced Topic: Sparse Matrix

- A sparse matrix is a matrix populated primarily with zeros.
- Save sparse matrix:
  - ✱ Dictionary of keys
  - ✱ List of lists
  - ✱ Coordinate list
  - ✱ Yale format
  - ✱ Etc.





# More Advanced Topic: Application using LA

- PageRank: imaging incredible large matrix
- Modern Digital imaging.
  - ✱ Video tracking: Xbox Kinect
- Genetics
- Cryptography
- Economic
- More ...



# More Advanced Topic: Application using LA

- Schedule & auto tuning
  - ✱ Test cases and pre-determined
  - ✱ Dynamically schedule
- Kernel and Convolution
  - ✱ Performs in parallel computers, edges of each blocks need to fix, according to the size of the kernel